

Coding for Reproducible Research

R Training for NISR

Luiza Andrade, Leonardo Viotti & Rob Marty

Aug. 2018



- 1 Introduction
- 2 Commenting
- 3 Creating a document outline in RStudio
- 4 Initial Settings
- 5 Using packages
- 6 Indenting
- 7 Functions inception
- 8 Section switches

Outline

- 1 Introduction
- 2 Commenting
- 3 Creating a document outline in RStudio
- 4 Initial Settings
- 5 Using packages
- 6 Indenting
- 7 Functions inception
- 8 Section switches
- 9 File paths
- 10 Looping

Objective

Create an R Master Script.

Content

- 1 Comments
- 2 Using R Studio to create a code index
- 3 Indenting
- 4 File paths
- 5 If statements
- 6 Using functions within functions
- 7 Packages
- 8 Loops

- The exercises on this section will help us create a Master R script
- To do them, go to the Code folder and open the file called *Lab 2 - Coding for Reproducible Research.R*
- A master script It's not specific to R, it's just good programming practice. Here we'll use it as an example to show you important features of R programming.

- A Master script is the easiest way to guarantee the replicability of your data work
- It has three main functions:
 - 1 Run all the scripts for your project
 - Without a master script, you either need to have one extremely long script, or write instructions about the order in which they should be run

- ② Create a road map to all data work
 - You should be able to reproduce all your work from raw data to all outputs by just running the master
 - Anyone should be able to follow and understand your work by reading the master

③ Allow easy collaboration

- If we share a project over DropBox or OneDrive all team member have the same folder structure
- A master script allows multiple people to set their own file paths to the project folder
- This way anyone sharing the project folder can easily run your codes

Outline

- 1 Introduction
- 2 **Commenting**
- 3 Creating a document outline in RStudio
- 4 Initial Settings
- 5 Using packages
- 6 Indenting
- 7 Functions inception
- 8 Section switches
- 9 File paths
- 10 Looping

- Comments is text that R will ignore when running your code
- Comments are the difference between instructions that are easy to follow or impossible to understand
- Comments are used to document two things:
 - ① What is being done in a given section of the code
 - ② Why it is being done
- Number 2 is what makes the difference between a well-commented code and a code that is just commented

- Let's take a look at the script we just opened
- You can see that the first few lines in the script are the header, but they're not commented out
- In R, errors will not always break your code, so you should still be able to run this script
- However, not commenting out comments is still bad practice, as it makes the code harder to read

- To comment a line, write `#` as its first character
- You can also add `#` half way through a line to comment whatever comes after it
- In Stata, you can use `/*` and `*/` to comment part of a line's code. That is not possible in R: whatever comes after `#` will be a comment
- To comment a selection of lines, press `Ctrl + Shift + C`

Exercise 1

Use the keyboard shortcut to comment the header of the script.

Exercise 2

Use the keyboard shortcut to comment the header of the script again.
What happened?

Outline

- 1 Introduction
- 2 Commenting
- 3 Creating a document outline in RStudio
- 4 Initial Settings
- 5 Using packages
- 6 Indenting
- 7 Functions inception
- 8 Section switches
- 9 File paths
- 10 Looping

Creating a document outline in RStudio

- RStudio also allows you to create an interactive index for your scripts
- To add a section to your code, create a commented line with the title of your section and add at least 4 trailing dashes, pound signs or equal signs after it

Exercise 3

Open the script index and make PART 0 a section header. Do the same for PART 1.

Exercise 4

Note that once you create a section header, an arrow appears right next to it. Click on the arrows of parts 0 and 1 to see what happens.

Creating a document outline in RStudio

- The outline can be accessed by clicking on the button on the top right corner of the script window. You can use it to jump from one section to another
- You can also use the keyboard shortcuts `Alt + L` (`Cmd + Option + L` on Mac) and `Alt + Shift + L` to collapse and expand sections

Outline

- 1 Introduction
- 2 Commenting
- 3 Creating a document outline in RStudio
- 4 Initial Settings**
- 5 Using packages
- 6 Indenting
- 7 Functions inception
- 8 Section switches
- 9 File paths
- 10 Looping

- A Stata do-file typically starts with a few settings:

```
clear  
set maxvar 120000  
set more off
```

Initial Settings

- We don't need to set the memory or the maximum number of variables in R, and more is automatically selected
- However, if you saved the last RStudio session in `.Rhistory`, the objects that were in RStudio's memory last time you closed it will still be there whenever you open it again
- Therefore, it's good practice to always clean the memory when starting a new session
- You can see all the objects currently in you memory in the *Environment* pane

Exercise 5

- 1 Make sure the *Environment* window is open (it should be empty now)
- 2 Create an object called `foo` with any content you pick
- 3 Type `ls()` to print the names of the object in memory
- 4 Type `rm(foo)` to remove it
- 5 To remove all objects, use `rm(list=ls())`

Outline

- 1 Introduction
- 2 Commenting
- 3 Creating a document outline in RStudio
- 4 Initial Settings
- 5 Using packages**
- 6 Indenting
- 7 Functions inception
- 8 Section switches
- 9 File paths
- 10 Looping

Using packages

- Since there is a lot of people developing for R, it can have many different functionalities
- To make it simpler, these functionalities are bundled into packages
- A package is the fundamental unit of shareable code
- It may contain new functions, but also more complex functionalities, such as a Graphic User Interface (GUI) or settings for parallel processing (similar to Stata MP)
- They can be shared through R's official repository - CRAN (10,000+ packages reviewed and tested) and many other online sources
- There are many other online sources such as Github, but it's important to be careful, as these probably haven't gone through a review process as rigorous as those in CRAN

Using packages

- To install and use packages you can either do it with the user interface or by the command prompt.

```
# Installing a package  
install.packages("stargazer",  
                 dependencies = T)  
# the dependencies argument also installs all other packages  
# that it may depend upon
```

- You only have to install a package once, but you have to load it every new session. To load a package type:

```
library(stargazer)
```

Using packages

Once a package is loaded, you can use its features and functions. Here's a list of some useful and cool packages:

- `Rcmdr`: Easy to use GUI
- `swirl`: An interactive learning environment for R and statistics.
- `ggplot2`: beautiful and versatile graphics (the syntax is a pain, though)
- `stargazer`: awesome latex regression and summary statistics tables
- `foreign`: reads dtas and other formats from other statistical software
- `zoo`: time series and panel data manipulation useful functions
- `data.table`: some functions to deal with huge data sets
- `sp` and `rgeos`: spatial analysis
- `multiwayvcov` and `sandwich`: clustered and robust standard errors
- `RODBC`, `RMySQL`, `RPostgresSQL`, `RSQLite`: for relational databases and using SQL in R.

Exercise 6

Install the `swirl` and `stargazer` packages, including packages necessary for them to run.

- TIP: use the helpfile to `install.packages` by typing `?install.packages` if you're not sure about how to do this.

Using packages

```
# Install stargazer and swirl  
install.packages(c("stargazer", "swirl"),  
                 dependencies = TRUE)
```

Exercise 7

Call the packages you just installed. Note that the `library` function only accepts one argument, so you will need to load each of them separately.

Using packages

```
library(stargazer)
```

```
##
```

```
## Please cite as:
```

```
## Hlavac, Marek (2018). stargazer: Well-Formatted Regression and Summary Statistics
```

```
## R package version 5.2.2. https://CRAN.R-project.org/package=stargazer
```

```
library(swirl)
```

```
##
```

```
## | Hi! Type swirl() when you are ready to begin.
```

Outline

- 1 Introduction
- 2 Commenting
- 3 Creating a document outline in RStudio
- 4 Initial Settings
- 5 Using packages
- 6 Indenting**
- 7 Functions inception
- 8 Section switches
- 9 File paths
- 10 Looping

- R does not distinguish between one empty space and many empty spaces, or one line break or many line breaks
- However, white space makes a big difference to the human eye and we would never share a Word document, an Excel sheet or a PowerPoint presentation without thinking about white space - although we call it formatting

Indenting

```
gen NoPlotDataBL = 0
replace NoPlotDataBL = 1 if c_plots_total_area >= .

gen NoHarvValueDataBL = 0
replace NoHarvValueDataBL = 1 if c_harv_value >= .

rename c_gross_yield c1_gross_yield
rename c_net_yield c1_net_yield
rename c_harv_value c1_harv_value
rename c_total_earnings c1_total_earnings
rename c_input_spend c2_inp_total_spending
rename c_IAAP_harv_value c1_IAAP_harv_value
rename c_plots_total_area c1_total_plotsize
rename c1_cropPlotShare_??? c1_cropPlotShare_all_???

tempfile BL_append
save `BL_append'
```

```
gen      NoPlotDataBL = 0
replace NoPlotDataBL = 1      if c_plots_total_area >= .

gen      NoHarvValueDataBL = 0
replace NoHarvValueDataBL = 1  if c_harv_value >= .

rename c_gross_yield      c1_gross_yield
rename c_net_yield        c1_net_yield
rename c_harv_value       c1_harv_value
rename c_total_earnings   c1_total_earnings
rename c_input_spend      c2_inp_total_spending
rename c_IAAP_harv_value  c1_IAAP_harv_value
rename c_plots_total_area c1_total_plotsize

rename c1_cropPlotShare_??? |c1_cropPlotShare_all_???

tempfile BL_append
save `BL_append'
```

- Indenting in R can be pretty different from what it looks like in Stata
- To indent a whole line, you can select that line and press Tab
- To unindent a whole line, you can select that line and press Shift + Tab
- However, this will not always work for different parts of a code in the same line
- In R, we typically don't introduce white space manually
- It's rather introduced by RStudio for us

Exercise 8

To see an example of how indenting works in RStudio, add a line between the two arguments of the `install.packages` function (the vector of package names and the `dependents` option). Then add a line between the two package names. Note that RStudio formats the different arguments of the function differently.

Indenting

```
# Same code, but this time easier to read  
install.packages(c("stargazer",  
                  "swirl"),  
                dependencies = TRUE)
```

Indenting

```
1 # Load panel data
2 panel<-read.csv(file.path(rawData,"lwh_panel.csv"))
3 # Create panel ID
4 panel$id<-(panel$hh_code*10000)+panel$year
5 # Check properties
6 sum(duplicated(panel$id))
7 # Subset data set
8 lwh<-panel[,c(id_vars,demographic_vars,yield_vars)]
9 # Turn numeric variable into factor
10 lwh$gender_hhh<-factor(lwh$gender_hhh,levels=c(0, 1),labels=c("Female","Male"))
```

```
1 # Load panel data
2 panel <- read.csv(file.path(rawData, "lwh_panel.csv"))
3
4 # Create panel ID
5 panel$id <- (panel$hh_code * 10000) + panel$year
6
7 # Check properties
8 sum(duplicated(panel$id))
9
10 # Subset data set
11 lwh <- panel[, c(id_vars,
12                 demographic_vars,
13                 yield_vars)]
14
15 # Turn numeric variable into factor
16 lwh$gender_hhh <- factor(lwh$gender_hhh,
17                          levels = c(0, 1),
18                          labels = c("Female", "Male"))
```

Outline

- 1 Introduction
- 2 Commenting
- 3 Creating a document outline in RStudio
- 4 Initial Settings
- 5 Using packages
- 6 Indenting
- 7 Functions inception**
- 8 Section switches
- 9 File paths
- 10 Looping

Functions inception

- In R, you can use the output of one function as the input of another, as long as they have the same format
- In fact, that's exactly what we just did when installing the packages
- To see that, select just the first argument of the `install.packages` function and press `Ctrl + Enter`
- The `c()` function, as we know, creates a vector with its arguments

```
c("stargazer", "swirl")
```

```
## [1] "stargazer" "swirl"
```

Functions inception

- The resulting vector is used as an input to the `install.packages` function
- We could also have stored this vector in the memory as an object and used that object as the input
- In fact, that's exactly what we are going to do next, so the code doesn't get too polluted as we add new packages

Functions inception

```
# Create packages object
packages <- c("stargazer",
              "swirl")

# Use it as an input to the install.packageS() function
install.packages(packages,
                 dependencies = TRUE)
```

Outline

- 1 Introduction
- 2 Commenting
- 3 Creating a document outline in RStudio
- 4 Initial Settings
- 5 Using packages
- 6 Indenting
- 7 Functions inception
- 8 Section switches**
- 9 File paths
- 10 Looping

Section switches

- Now, installing packages can be time-consuming, especially as the number of packages grow, and each package only needs to be installed once
- Adding switches to select what parts of the code to run allows you to only install the packages when you're using a computer that doesn't have them installed yet
- Adding switches is more efficient than commenting parts of the code out, because you can see all switches before running the code, so that avoids the mistake of saving a code with a commented section and forgetting to uncomment it later
- In Stata, section switches would be saved as locals. In R, the equivalent to that would be to create a new object

If statements

- To add a switch, you first create a dummy object with a self-explanatory name
- Then, you create an if statement that runs the code if that switch is on
- If statements in R look like this:

```
# Turn switch on  
PRINT_NAME <- 1  
  
# Install packages  
if (PRINT_NAME == 1) {  
  print(Sys.getenv("USERNAME"))  
}
```

Exercise 10

Create a switch called `INSTALL_PACKAGES` and an if statement that only runs the `install.packages` function if the switch is activated.

- TIP: Section switches can also be Boolean objects.

Section switches

```
# Turn switch on
INSTALL_PACKAGES <- 1

# Install packages
if (INSTALL_PACKAGES == 1) {
  install.packages(packages,
                  dependencies = TRUE)
}
```

Section switches

- Possible variations would include

```
# Turn switch on
```

```
INSTALL_PACKAGES <- TRUE
```

```
# Using a Boolean object
```

```
if (INSTALL_PACKAGES == TRUE) {  
  install.packages/packages, dep = T)  
}
```

```
# Which is the same as
```

```
if (INSTALL_PACKAGES) {  
  install.packages/packages, dep = T)  
}
```

Outline

- 1 Introduction
- 2 Commenting
- 3 Creating a document outline in RStudio
- 4 Initial Settings
- 5 Using packages
- 6 Indenting
- 7 Functions inception
- 8 Section switches
- 9 File paths**
- 10 Looping

- The next important part of a Master script are file paths
- We suggest always using **explicit** and **dynamic** file paths

- Implicit and static file path:

```
# Set working directory  
setwd("C:/Users/luiza/Documents/GitHub/R-Training/  
      DataWork/DataSets/Final")  
  
# Load data set  
read.csv("lwh_clean.csv",  
         header = T)
```


- Explicit and static file path:

```
# Load data set  
read.csv("C:/Users/luiza/Documents/GitHub/R-Training/  
DataWork/DataSets/Final/lwh_clean.csv",  
         header = T)
```

- Explicit and dynamic file path:

```
# Define dynamic file path
finalData <- "C:/Users/luiza/Documents/GitHub/R-Training/
             DataWork/DataSets/Final"

# Load data set
lwh <- read.csv(file.path(finalData,"lwh_clean.csv"),
                header = T)
```

- Using dynamic file paths makes collaboration easier, since every user only needs to add their folder path once before running all the codes
- Using explicit file paths prevents mistakes. For example, when running just a few lines of code instead of the whole script, it's common to forget to run the line setting the directory and have output saved in the wrong folder

- File paths in R, as in Stata, are basically just strings
- Note, however, that in R we can only use forward slashes (/) to separate folder names

Exercise 11

Let's start by adding the folder path to the training's folder in your computer to the beginning of PART 3.

- You can set file paths in your master using the `file.path()` function
- This function concatenates strings using `/` as a separator to create file paths

File paths

```
# Project folder
projectFolder <-
  "C:/Users/luiza/Documents/GitHub/R-Training"

# Data work folder
dataWorkFolder <- file.path(projectFolder, "DataWork")

# Print data work folder
dataWorkFolder
```

```
## [1] "C:/Users/luiza/Documents/GitHub/R-Training/DataWork"
```

Let's check if that worked, as we will need your Master script to be running smoothly for the other sessions.

Exercise 12

- 1 Turn off the switch that installs packages in your Master script
- 2 Run the whole script
- 3 Type the following code to open a data set using the file paths you just set:

```
# Load data set  
lwh <- read.csv(file.path(rawData, "lwh_panel.csv"),  
                header = T)
```

Outline

- 1 Introduction
- 2 Commenting
- 3 Creating a document outline in RStudio
- 4 Initial Settings
- 5 Using packages
- 6 Indenting
- 7 Functions inception
- 8 Section switches
- 9 File paths
- 10 Looping

- We're almost at the end of this section
- But the DRY rule can still be applied to part of this code

The DRY rule:

DONT REPEAT YOURSELF

Don't repeat yourself

- This is a rule borrow from computer science
- When coding, we often have to repeat the same operation multiple times
- If you do this by just copying and pasting a piece of code and changing a few arguments, it's easy to make mistakes such as forgetting to change the argument once or changing the wrong argument
- Copying and pasting a piece of code multiple times can also make your code really long and difficult to read
- Creating a loop may take more time to set up, but it's easier to read and reduces mistakes
- In particular, fixing bugs and adjusting the code is much quicker, since you only need to do it once

Looping

- In Stata, we'd usually use a `foreach` loop to go through a list of objects
- The equivalent to that in R would be to write a `for` loop like this

```
# A for loop in R  
for (number in c(1.2,2.5)) {  
  print(round(number))  
}
```

```
## [1] 1
```

```
## [1] 2
```

- R, however, has a whole function family that allows users to loop through an object in a more efficient way
- They're called `apply` and there are many of them, for different use cases.
- For the purpose of this training, we will only use two of them, `sapply` and `apply`
- If you look for the `apply` help file, you can see all of them

Looping

- `sapply(X, FUN, ...)`: applies a function to all elements of a vector or list and returns the result in a vector. Its arguments are
 - **X**: a matrix (or data frame) the function will be applied to
 - **FUN**: the function you want to apply
 - **...**: possible function options

```
# A much more elegant for loop in R  
sapply(c(1.2,2.5), round)
```

```
## [1] 1 2
```

Exercise 13

Use the `sapply()` function to apply the `library()` function to all packages you have selected. TIP: Set the `character.only` argument equal to `TRUE`

```
# Load all listed packages  
sapply(packages,  
       library, character.only = TRUE)  
# it's necessary to use the character.only option because  
# originally the library function did not accept strings
```

A more generic version is the `apply` function.

- `apply(X, MARGIN, FUN, ...)`: applies a function to all columns or rows of matrix. Its arguments are
 - **X**: a matrix (or data frame) the function will be applied to
 - **MARGIN**: 1 to apply the function to all rows or 2 to apply the function to all columns
 - **FUN**: the function you want to apply
 - **...**: possible function options

Looping

```
# Create a matrix  
matrix <- matrix(c(1, 24, 9, 6, 9, 4, 2, 74, 2),  
                 nrow = 3)
```

```
# Look at the matrix  
matrix
```

```
##      [,1] [,2] [,3]  
## [1,]    1    6    2  
## [2,]   24    9   74  
## [3,]    9    4    2
```

Looping

```
# Row means  
apply(matrix, 1, mean)
```

```
## [1] 3.00000 35.66667 5.00000
```

```
# Column means  
apply(matrix, 2, mean)
```

```
## [1] 11.333333 6.333333 26.000000
```

That's all, folks

- Now you have a template master script to use across this training's sessions
- Save the script that you created during this session in the *DataWork* folder. Call it *MASTER.R*
- You can run scripts from the Master script by using the `source()` function as you write scripts for future sessions, but we will do that on the next session

That's all, folks

Homework

In the next slide here's a list of all the packages we'll need for the next sessions. Installing them might take a while, so paste them to your Master script and run it again before the next session:

```
packages <- c("readstata13", "foreign",  
              "doBy", "broom",  
              "stargazer",  
              "ggplot2", "plotly", "ggrepel",  
              "RColorBrewer",  
              "sp", "rgdal", "rgeos", "raster", "velox",  
              "ggmap", "rasterVis", "leaflet")
```

Thank you!

Bonus exercise

- You can customize your loops in R by defining your own function
- This is done using a function conveniently called `function()`
- For example, if instead of just printing a number we want to print it's square, we could create a function that does both:

```
# A much more elegant for loop in R  
sapply(c(1,2), function(x) x^2)
```

```
## [1] 1 4
```